

NAME

environ, execl, execl, execlp, execv, execve, execvp, fexecve — execute a file

SYNOPSIS

```
#include <unistd.h>

extern char **environ;
int execl(const char *path, const char *arg0, ... /*, (char *)0 */);
int execl(const char *path, const char *arg0, ... /*,
        (char *)0, char *const envp[] */);
int execlp(const char *file, const char *arg0, ... /*, (char *)0 */);
int execv(const char *path, char *const argv[]);
int execve(const char *path, char *const argv[], char *const envp[]);
int execvp(const char *file, char *const argv[]);
int fexecve(int fd, char *const argv[], char *const envp[]);
```

DESCRIPTION

The *exec* family of functions shall replace the current process image with a new process image. The new image shall be constructed from a regular, executable file called the *new process image file*. There shall be no return from a successful *exec*, because the calling process image is overlaid by the new process image.

The *fexecve()* function shall be equivalent to the *execve()* function except that the file to be executed is determined by the file descriptor *fd* instead of a pathname. The file offset of *fd* is ignored.

When a C-language program is executed as a result of a call to one of the *exec* family of functions, it shall be entered as a C-language function call as follows:

```
int main (int argc, char *argv[]);
```

where *argc* is the argument count and *argv* is an array of character pointers to the arguments themselves. In addition, the following variable:

```
extern char **environ;
```

is initialized as a pointer to an array of character pointers to the environment strings. The *argv* and *environ* arrays are each terminated by a null pointer. The null pointer terminating the *argv* array is not counted in *argc*.

Conforming multi-threaded applications shall not use the *environ* variable to access or modify any environment variable while any other thread is concurrently modifying any environment variable. A call to any function dependent on any environment variable shall be considered a use of the *environ* variable to access that environment variable.

The arguments specified by a program with one of the *exec* functions shall be passed on to the new process image in the corresponding *main()* arguments.

The argument *path* points to a pathname that identifies the new process image file.

The argument *file* is used to construct a pathname that identifies the new process image file. If the *file* argument contains a <slash> character, the *file* argument shall be used as the pathname for this file. Otherwise, the path prefix for this file is obtained by a search of the directories passed as the environment variable *PATH* (see XBD [Chapter 8](#), on page 173). If this environment variable is not present, the results of the search are implementation-defined.

There are two distinct ways in which the contents of the process image file may cause the execution to fail, distinguished by the setting of *errno* to either [ENOEXEC] or [EINVAL] (see the **ERRORS** section). In the cases where the other members of the *exec* family of functions would

fail and set *errno* to [ENOEXEC], the *execlp()* and *execvp()* functions shall execute a command interpreter and the environment of the executed command shall be as if the process invoked the *sh* utility using *execl()* as follows:

```
execl(<shell path>, arg0, file, arg1, ..., (char *)0);
```

where *<shell path>* is an unspecified pathname for the *sh* utility, *file* is the process image file, and for *execvp()*, where *arg0*, *arg1*, and so on correspond to the values passed to *execvp()* in *argv[0]*, *argv[1]*, and so on.

The arguments represented by *arg0*,... are pointers to null-terminated character strings. These strings shall constitute the argument list available to the new process image. The list is terminated by a null pointer. The argument *arg0* should point to a filename that is associated with the process being started by one of the *exec* functions.

The argument *argv* is an array of character pointers to null-terminated strings. The application shall ensure that the last member of this array is a null pointer. These strings shall constitute the argument list available to the new process image. The value in *argv[0]* should point to a filename that is associated with the process being started by one of the *exec* functions.

The argument *envp* is an array of character pointers to null-terminated strings. These strings shall constitute the environment for the new process image. The *envp* array is terminated by a null pointer.

For those forms not containing an *envp* pointer (*execl()*, *execv()*, *execlp()*, and *execvp()*), the environment for the new process image shall be taken from the external variable *environ* in the calling process.

The number of bytes available for the new process' combined argument and environment lists is {ARG_MAX}. It is implementation-defined whether null terminators, pointers, and/or any alignment bytes are included in this total.

File descriptors open in the calling process image shall remain open in the new process image, except for those whose close-on-exec flag FD_CLOEXEC is set. For those file descriptors that remain open, all attributes of the open file description remain unchanged. For any file descriptor that is closed for this reason, file locks are removed as a result of the close as described in *close()*. Locks that are not removed by closing of file descriptors remain unchanged.

If file descriptors 0, 1, and 2 would otherwise be closed after a successful call to one of the *exec* family of functions, and the new process image file has the set-user-ID or set-group-ID file mode bits set, and the ST_NOSUID bit is not set for the file system containing the new process image file, implementations may open an unspecified file for each of these file descriptors in the new process image.

Directory streams open in the calling process image shall be closed in the new process image.

The state of the floating-point environment in the initial thread of the new process image shall be set to the default.

The state of conversion descriptors and message catalog descriptors in the new process image is undefined.

For the new process image, the equivalent of:

```
setlocale(LC_ALL, "C")
```

shall be executed at start-up.

Signals set to the default action (SIG_DFL) in the calling process image shall be set to the default action in the new process image. Except for SIGCHLD, signals set to be ignored (SIG_IGN) by

25772		the calling process image shall be set to be ignored by the new process image. Signals set to be
25773		caught by the calling process image shall be set to the default action in the new process image
25774		(see <signal.h>).
25775		If the SIGCHLD signal is set to be ignored by the calling process image, it is unspecified whether
25776		the SIGCHLD signal is set to be ignored or to the default action in the new process image.
25777	XSI	After a successful call to any of the <i>exec</i> functions, alternate signal stacks are not preserved and
25778		the SA_ONSTACK flag shall be cleared for all signals.
25779		After a successful call to any of the <i>exec</i> functions, any functions previously registered by the
25780		<i>atexit()</i> or <i>pthread_atfork()</i> functions are no longer registered.
25781	XSI	If the ST_NOSUID bit is set for the file system containing the new process image file, then the
25782		effective user ID, effective group ID, saved set-user-ID, and saved set-group-ID are unchanged
25783		in the new process image. Otherwise, if the set-user-ID mode bit of the new process image file is
25784		set, the effective user ID of the new process image shall be set to the user ID of the new process
25785		image file. Similarly, if the set-group-ID mode bit of the new process image file is set, the
25786		effective group ID of the new process image shall be set to the group ID of the new process
25787		image file. The real user ID, real group ID, and supplementary group IDs of the new process
25788		image shall remain the same as those of the calling process image. The effective user ID and
25789		effective group ID of the new process image shall be saved (as the saved set-user-ID and the
25790		saved set-group-ID) for use by <i>setuid()</i> .
25791	XSI	Any shared memory segments attached to the calling process image shall not be attached to the
25792		new process image.
25793		Any named semaphores open in the calling process shall be closed as if by appropriate calls to
25794		<i>sem_close()</i> .
25795	TYM	Any blocks of typed memory that were mapped in the calling process are unmapped, as if
25796		<i>munmap()</i> was implicitly called to unmap them.
25797	ML	Memory locks established by the calling process via calls to <i>mlockall()</i> or <i>mlock()</i> shall be
25798		removed. If locked pages in the address space of the calling process are also mapped into the
25799		address spaces of other processes and are locked by those processes, the locks established by the
25800		other processes shall be unaffected by the call by this process to the <i>exec</i> function. If the <i>exec</i>
25801		function fails, the effect on memory locks is unspecified.
25802		Memory mappings created in the process are unmapped before the address space is rebuilt for
25803		the new process image.
25804	SS	When the calling process image does not use the SCHED_FIFO, SCHED_RR, or
25805		SCHED_SPORADIC scheduling policies, the scheduling policy and parameters of the new
25806		process image and the initial thread in that new process image are implementation-defined.
25807	PS	When the calling process image uses the SCHED_FIFO, SCHED_RR, or SCHED_SPORADIC
25808		scheduling policies, the process policy and scheduling parameter settings shall not be changed
25809	TPS	by a call to an <i>exec</i> function. The initial thread in the new process image shall inherit the process
25810		scheduling policy and parameters. It shall have the default system contention scope, but shall
25811		inherit its allocation domain from the calling process image.
25812		Per-process timers created by the calling process shall be deleted before replacing the current
25813		process image with the new process image.
25814	MSG	All open message queue descriptors in the calling process shall be closed, as described in
25815		<i>mq_close()</i> .
25816		Any outstanding asynchronous I/O operations may be canceled. Those asynchronous I/O

25817		operations that are not canceled shall complete as if the <i>exec</i> function had not yet occurred, but
25818		any associated signal notifications shall be suppressed. It is unspecified whether the <i>exec</i>
25819		function itself blocks awaiting such I/O completion. In no event, however, shall the new process
25820		image created by the <i>exec</i> function be affected by the presence of outstanding asynchronous I/O
25821		operations at the time the <i>exec</i> function is called. Whether any I/O is canceled, and which I/O
25822		may be canceled upon <i>exec</i> , is implementation-defined.
25823	CPT	The new process image shall inherit the CPU-time clock of the calling process image. This
25824		inheritance means that the process CPU-time clock of the process being <i>exec</i> -ed shall not be
25825		reinitialized or altered as a result of the <i>exec</i> function other than to reflect the time spent by the
25826		process executing the <i>exec</i> function itself.
25827	TCT	The initial value of the CPU-time clock of the initial thread of the new process image shall be set
25828		to zero.
25829	OB TRC	If the calling process is being traced, the new process image shall continue to be traced into the
25830		same trace stream as the original process image, but the new process image shall not inherit the
25831		mapping of trace event names to trace event type identifiers that was defined by calls to the
25832		<i>posix_trace_eventid_open()</i> or the <i>posix_trace_trid_eventid_open()</i> functions in the calling process
25833		image.
25834		If the calling process is a trace controller process, any trace streams that were created by the
25835		calling process shall be shut down as described in the <i>posix_trace_shutdown()</i> function.
25836		The thread ID of the initial thread in the new process image is unspecified.
25837		The size and location of the stack on which the initial thread in the new process image runs is
25838		unspecified.
25839		The initial thread in the new process image shall have its cancellation type set to
25840		PTHREAD_CANCEL_DEFERRED and its cancellation state set to
25841		PTHREAD_CANCEL_ENABLED.
25842		The initial thread in the new process image shall have all thread-specific data values set to
25843		NULL and all thread-specific data keys shall be removed by the call to <i>exec</i> without running
25844		destructors.
25845		The initial thread in the new process image shall be joinable, as if created with the <i>detachstate</i>
25846		attribute set to PTHREAD_CREATE_JOINABLE.
25847		The new process shall inherit at least the following attributes from the calling process image:
25848	XSI	• Nice value (see <i>nice()</i>)
25849	XSI	• <i>semadj</i> values (see <i>semop()</i>)
25850		• Process ID
25851		• Parent process ID
25852		• Process group ID
25853		• Session membership
25854		• Real user ID
25855		• Real group ID
25856		• Supplementary group IDs

25857		• Time left until an alarm clock signal (see <i>alarm()</i>)
25858		• Current working directory
25859		• Root directory
25860		• File mode creation mask (see <i>umask()</i>)
25861	XSI	• File size limit (see <i>getrlimit()</i> and <i>setrlimit()</i>)
25862		• Process signal mask (see <i>pthread_sigmask()</i>)
25863		• Pending signal (see <i>sigpending()</i>)
25864		• <i>tms_utime</i> , <i>tms_stime</i> , <i>tms_cutime</i> , and <i>tms_cstime</i> (see <i>times()</i>)
25865	XSI	• Resource limits
25866		• Controlling terminal
25867	XSI	• Interval timers
25868		The initial thread of the new process shall inherit at least the following attributes from the
25869		calling thread:
25870		• Signal mask (see <i>sigprocmask()</i> and <i>pthread_sigmask()</i>)
25871		• Pending signals (see <i>sigpending()</i>)
25872		All other process attributes defined in this volume of POSIX.1-2008 shall be inherited in the new
25873		process image from the old process image. All other thread attributes defined in this volume of
25874		POSIX.1-2008 shall be inherited in the initial thread in the new process image from the calling
25875		thread in the old process image. The inheritance of process or thread attributes not defined by
25876		this volume of POSIX.1-2008 is implementation-defined.
25877		A call to any <i>exec</i> function from a process with more than one thread shall result in all threads
25878		being terminated and the new executable image being loaded and executed. No destructor
25879		functions or cleanup handlers shall be called.
25880		Upon successful completion, the <i>exec</i> functions shall mark for update the last data access
25881		timestamp of the file. If an <i>exec</i> function failed but was able to locate the process image file,
25882		whether the last data access timestamp is marked for update is unspecified. Should the <i>exec</i>
25883		function succeed, the process image file shall be considered to have been opened with <i>open()</i> .
25884		The corresponding <i>close()</i> shall be considered to occur at a time after this open, but before
25885		process termination or successful completion of a subsequent call to one of the <i>exec</i> functions,
25886		<i>posix_spawn()</i> , or <i>posix_spawnnp()</i> . The <i>argv[]</i> and <i>envp[]</i> arrays of pointers and the strings to
25887		which those arrays point shall not be modified by a call to one of the <i>exec</i> functions, except as a
25888		consequence of replacing the process image.
25889	XSI	The saved resource limits in the new process image are set to be a copy of the process'
25890		corresponding hard and soft limits.
25891		RETURN VALUE
25892		If one of the <i>exec</i> functions returns to the calling process image, an error has occurred; the return
25893		value shall be <i>-1</i> , and <i>errno</i> shall be set to indicate the error.
25894		ERRORS
25895		The <i>exec</i> functions shall fail if:
25896	[E2BIG]	The number of bytes used by the new process image's argument list and
25897		environment list is greater than the system-imposed limit of {ARG_MAX}
25898		bytes.

25899	[EACCES]	Search permission is denied for a directory listed in the new process image file's path prefix, or the new process image file denies execution permission, or the new process image file is not a regular file and the implementation does not support execution of files of its type.
25900		
25901		
25902		
25903	[EINVAL]	The new process image file has appropriate privileges and has a recognized executable binary format, but the system does not support execution of a file with this format.
25904		
25905		
25906	[ELOOP]	A loop exists in symbolic links encountered during resolution of the <i>path</i> or <i>file</i> argument.
25907		
25908	[ENAMETOOLONG]	
25909		The length of a component of a pathname is longer than {NAME_MAX}.
25910	[ENOENT]	A component of <i>path</i> or <i>file</i> does not name an existing file or <i>path</i> or <i>file</i> is an empty string.
25911		
25912	[ENOTDIR]	A component of the new process image file's path prefix is not a directory, or the new process image file's pathname contains at least one non- <i><slash></i> character and ends with one or more trailing <i><slash></i> characters and the last pathname component names an existing file that is neither a directory nor a symbolic link to a directory.
25913		
25914		
25915		
25916		
25917		The <i>exec</i> functions, except for <i>execlp()</i> and <i>execvp()</i> , shall fail if:
25918	[ENOEXEC]	The new process image file has the appropriate access permission but has an unrecognized format.
25919		
25920		The <i>fexecve()</i> function shall fail if:
25921	[EBADF]	The <i>fd</i> argument is not a valid file descriptor open for executing.
25922		The <i>exec</i> functions may fail if:
25923	[ELOOP]	More than {SYMLOOP_MAX} symbolic links were encountered during resolution of the <i>path</i> or <i>file</i> argument.
25924		
25925	[ENAMETOOLONG]	
25926		The length of the <i>path</i> argument or the length of the pathname constructed from the <i>file</i> argument exceeds {PATH_MAX}, or pathname resolution of a symbolic link produced an intermediate result with a length that exceeds {PATH_MAX}.
25927		
25928		
25929		
25930	[ENOMEM]	The new process image requires more memory than is allowed by the hardware or system-imposed memory management constraints.
25931		
25932	[ETXTBSY]	The new process image file is a pure procedure (shared text) file that is currently open for writing by some process.
25933		

25934 **EXAMPLES**25935 **Using execl()**

25936 The following example executes the *ls* command, specifying the pathname of the executable
25937 (*/bin/ls*) and using arguments supplied directly to the command to produce single-column
25938 output.

```
25939 #include <unistd.h>
25940 int ret;
25941 ...
25942 ret = execl ("/bin/ls", "ls", "-l", (char *)0);
```

25943 **Using execle()**

25944 The following example is similar to [Using execl\(\)](#). In addition, it specifies the environment for
25945 the new process image using the *env* argument.

```
25946 #include <unistd.h>
25947 int ret;
25948 char *env[] = { "HOME=/usr/home", "LOGNAME=home", (char *)0 };
25949 ...
25950 ret = execle ("/bin/ls", "ls", "-l", (char *)0, env);
```

25951 **Using execlp()**

25952 The following example searches for the location of the *ls* command among the directories
25953 specified by the *PATH* environment variable.

```
25954 #include <unistd.h>
25955 int ret;
25956 ...
25957 ret = execlp ("ls", "ls", "-l", (char *)0);
```

25958 **Using execv()**

25959 The following example passes arguments to the *ls* command in the *cmd* array.

```
25960 #include <unistd.h>
25961 int ret;
25962 char *cmd[] = { "ls", "-l", (char *)0 };
25963 ...
25964 ret = execv ("/bin/ls", cmd);
```

Using `execve()`

The following example passes arguments to the *ls* command in the *cmd* array, and specifies the environment for the new process image using the *env* argument.

```
#include <unistd.h>

int ret;
char *cmd[] = { "ls", "-l", (char *)0 };
char *env[] = { "HOME=/usr/home", "LOGNAME=home", (char *)0 };
...
ret = execve ("/bin/ls", cmd, env);
```

Using `execvp()`

The following example searches for the location of the *ls* command among the directories specified by the *PATH* environment variable, and passes arguments to the *ls* command in the *cmd* array.

```
#include <unistd.h>

int ret;
char *cmd[] = { "ls", "-l", (char *)0 };
...
ret = execvp ("ls", cmd);
```

APPLICATION USAGE

As the state of conversion descriptors and message catalog descriptors in the new process image is undefined, conforming applications should not rely on their use and should close them prior to calling one of the *exec* functions.

Applications that require other than the default POSIX locale should call *setlocale()* with the appropriate parameters to establish the locale of the new process.

The *environ* array should not be accessed directly by the application.

The new process might be invoked in a non-conforming environment if the *envp* array does not contain implementation-defined variables required by the implementation to provide a conforming environment. See the *_CS_V7_ENV* entry in **<unistd.h>** and *confstr()* for details.

Applications should not depend on file descriptors 0, 1, and 2 being closed after an *exec*. A future version may allow these file descriptors to be automatically opened for any process.

If an application wants to perform a checksum test of the file being executed before executing it, the file will need to be opened with read permission to perform the checksum test.

Since execute permission is checked by *fexecve()*, the file description *fd* need not have been opened with the *O_EXEC* flag. However, if the file to be executed denies read and write permission for the process preparing to do the *exec*, the only way to provide the *fd* to *fexecve()* will be to use the *O_EXEC* flag when opening *fd*. In this case, the application will not be able to perform a checksum test since it will not be able to read the contents of the file.

Note that when a file descriptor is opened with *O_RDONLY*, *O_RDWR*, or *O_WRONLY* mode, the file descriptor can be used to read, read and write, or write the file, respectively, even if the mode of the file changes after the file was opened. Using the *O_EXEC* open mode is different; *fexecve()* will ignore the mode that was used when the file descriptor was opened and the *exec* will fail if the mode of the file associated with *fd* does not grant execute permission to the calling process at the time *fexecve()* is called.

RATIONALE

Early proposals required that the value of *argc* passed to *main()* be “one or greater”. This was driven by the same requirement in drafts of the ISO C standard. In fact, historical implementations have passed a value of zero when no arguments are supplied to the caller of the *exec* functions. This requirement was removed from the ISO C standard and subsequently removed from this volume of POSIX.1-2008 as well. The wording, in particular the use of the word *should*, requires a Strictly Conforming POSIX Application to pass at least one argument to the *exec* function, thus guaranteeing that *argc* be one or greater when invoked by such an application. In fact, this is good practice, since many existing applications reference *argv*[0] without first checking the value of *argc*.

The requirement on a Strictly Conforming POSIX Application also states that the value passed as the first argument be a filename associated with the process being started. Although some existing applications pass a pathname rather than a filename in some circumstances, a filename is more generally useful, since the common usage of *argv*[0] is in printing diagnostics. In some cases the filename passed is not the actual filename of the file; for example, many implementations of the *login* utility use a convention of prefixing a <hyphen> (‘-’) to the actual filename, which indicates to the command interpreter being invoked that it is a “login shell”.

Historically, there have been two ways that implementations can *exec* shell scripts.

One common historical implementation is that the *execl()*, *execv()*, *execle()*, and *execve()* functions return an [ENOEXEC] error for any file not recognizable as executable, including a shell script. When the *execlp()* and *execvp()* functions encounter such a file, they assume the file to be a shell script and invoke a known command interpreter to interpret such files. This is now required by POSIX.1-2008. These implementations of *execvp()* and *execlp()* only give the [ENOEXEC] error in the rare case of a problem with the command interpreter’s executable file. Because of these implementations, the [ENOEXEC] error is not mentioned for *execlp()* or *execvp()*, although implementations can still give it.

Another way that some historical implementations handle shell scripts is by recognizing the first two bytes of the file as the character string “#!” and using the remainder of the first line of the file as the name of the command interpreter to execute.

One potential source of confusion noted by the standard developers is over how the contents of a process image file affect the behavior of the *exec* family of functions. The following is a description of the actions taken:

1. If the process image file is a valid executable (in a format that is executable and valid and having appropriate privileges) for this system, then the system executes the file.
2. If the process image file has appropriate privileges and is in a format that is executable but not valid for this system (such as a recognized binary for another architecture), then this is an error and *errno* is set to [EINVAL] (see later RATIONALE on [EINVAL]).
3. If the process image file has appropriate privileges but is not otherwise recognized:
 - a. If this is a call to *execlp()* or *execvp()*, then they invoke a command interpreter assuming that the process image file is a shell script.
 - b. If this is not a call to *execlp()* or *execvp()*, then an error occurs and *errno* is set to [ENOEXEC].

Applications that do not require to access their arguments may use the form:

```
main(void)
```

as specified in the ISO C standard. However, the implementation will always provide the two

arguments *argc* and *argv*, even if they are not used.

Some implementations provide a third argument to *main()* called *envp*. This is defined as a pointer to the environment. The ISO C standard specifies invoking *main()* with two arguments, so implementations must support applications written this way. Since this volume of POSIX.1-2008 defines the global variable *environ*, which is also provided by historical implementations and can be used anywhere that *envp* could be used, there is no functional need for the *envp* argument. Applications should use the *getenv()* function rather than accessing the environment directly via either *envp* or *environ*. Implementations are required to support the two-argument calling sequence, but this does not prohibit an implementation from supporting *envp* as an optional third argument.

This volume of POSIX.1-2008 specifies that signals set to SIG_IGN remain set to SIG_IGN, and that the new process image inherits the signal mask of the thread that called *exec* in the old process image. This is consistent with historical implementations, and it permits some useful functionality, such as the *nohup* command. However, it should be noted that many existing applications wrongly assume that they start with certain signals set to the default action and/or unblocked. In particular, applications written with a simpler signal model that does not include blocking of signals, such as the one in the ISO C standard, may not behave properly if invoked with some signals blocked. Therefore, it is best not to block or ignore signals across *execs* without explicit reason to do so, and especially not to block signals across *execs* of arbitrary (not closely co-operating) programs.

The *exec* functions always save the value of the effective user ID and effective group ID of the process at the completion of the *exec*, whether or not the set-user-ID or the set-group-ID bit of the process image file is set.

The statement about *argv[]* and *envp[]* being constants is included to make explicit to future writers of language bindings that these objects are completely constant. Due to a limitation of the ISO C standard, it is not possible to state that idea in standard C. Specifying two levels of *const-qualification* for the *argv[]* and *envp[]* parameters for the *exec* functions may seem to be the natural choice, given that these functions do not modify either the array of pointers or the characters to which the function points, but this would disallow existing correct code. Instead, only the array of pointers is noted as constant. The table of assignment compatibility for *dst=src* derived from the ISO C standard summarizes the compatibility:

<i>dst:</i>	char *[]	const char *[]	char *const[]	const char *const[]
<i>src:</i>				
char *[]	VALID	—	VALID	—
const char *[]	—	VALID	—	VALID
char * const []	—	—	VALID	—
const char *const[]	—	—	—	VALID

Since all existing code has a source type matching the first row, the column that gives the most valid combinations is the third column. The only other possibility is the fourth column, but using it would require a cast on the *argv* or *envp* arguments. It is unfortunate that the fourth column cannot be used, because the declaration a non-expert would naturally use would be that in the second row.

The ISO C standard and this volume of POSIX.1-2008 do not conflict on the use of *environ*, but some historical implementations of *environ* may cause a conflict. As long as *environ* is treated in the same way as an entry point (for example, *fork()*), it conforms to both standards. A library can contain *fork()*, but if there is a user-provided *fork()*, that *fork()* is given precedence and no problem ensues. The situation is similar for *environ*: the definition in this volume of POSIX.1-2008 is to be used if there is no user-provided *environ* to take precedence. At least three

26101	implementations are known to exist that solve this problem.
26102	[E2BIG] The limit {ARG_MAX} applies not just to the size of the argument list, but to
26103	the sum of that and the size of the environment list.
26104	[EFAULT] Some historical systems return [EFAULT] rather than [ENOEXEC] when the
26105	new process image file is corrupted. They are non-conforming.
26106	[EINVAL] This error condition was added to POSIX.1-2008 to allow an implementation
26107	to detect executable files generated for different architectures, and indicate this
26108	situation to the application. Historical implementations of shells, <i>execvp()</i> , and
26109	<i>execlp()</i> that encounter an [ENOEXEC] error will execute a shell on the
26110	assumption that the file is a shell script. This will not produce the desired
26111	effect when the file is a valid executable for a different architecture. An
26112	implementation may now choose to avoid this problem by returning
26113	[EINVAL] when a valid executable for a different architecture is encountered.
26114	Some historical implementations return [EINVAL] to indicate that the <i>path</i>
26115	argument contains a character with the high order bit set. The standard
26116	developers chose to deviate from historical practice for the following reasons:
26117	1. The new utilization of [EINVAL] will provide some measure of utility
26118	to the user community.
26119	2. Historical use of [EINVAL] is not acceptable in an internationalized
26120	operating environment.
26121	[ENAMETOOLONG]
26122	Since the file pathname may be constructed by taking elements in the <i>PATH</i>
26123	variable and putting them together with the filename, the
26124	[ENAMETOOLONG] error condition could also be reached this way.
26125	[ETXTBSY] System V returns this error when the executable file is currently open for
26126	writing by some process. This volume of POSIX.1-2008 neither requires nor
26127	prohibits this behavior.
26128	Other systems (such as System V) may return [EINTR] from <i>exec</i> . This is not addressed by this
26129	volume of POSIX.1-2008, but implementations may have a window between the call to <i>exec</i> and
26130	the time that a signal could cause one of the <i>exec</i> calls to return with [EINTR].
26131	An explicit statement regarding the floating-point environment (as defined in the <fenv.h>
26132	header) was added to make it clear that the floating-point environment is set to its default when
26133	a call to one of the <i>exec</i> functions succeeds. The requirements for inheritance or setting to the
26134	default for other process and thread start-up functions is covered by more generic statements in
26135	their descriptions and can be summarized as follows:
26136	<i>posix_spawn()</i> Set to default.
26137	<i>fork()</i> Inherit.
26138	<i>pthread_create()</i> Inherit.
26139	The purpose of the <i>fexecve()</i> function is to enable executing a file which has been verified to be
26140	the intended file. It is possible to actively check the file by reading from the file descriptor and be
26141	sure that the file is not exchanged for another between the reading and the execution.
26142	Alternatively, an function like <i>openat()</i> can be used to open a file which has been found by
26143	reading the content of a directory using <i>readdir()</i> .

26144 **FUTURE DIRECTIONS**

26145 None.

26146 **SEE ALSO**

26147 *alarm()*, *atexit()*, *chmod()*, *close()*, *confstr()*, *exit()*, *fcntl()*, *fork()*, *fstatvfs()*, *getenv()*, *getitimer()*,
 26148 *getrlimit()*, *mknod()*, *mmap()*, *nice()*, *open()*, *posix_spawn()*, *posix_trace_create()*,
 26149 *posix_trace_event()*, *posix_trace_eventid_equal()*, *pthread_atfork()*, *pthread_sigmask()*, *putenv()*,
 26150 *readdir()*, *semop()*, *setlocale()*, *shmat()*, *sigaction()*, *sigaltstack()*, *sigpending()*, *system()*, *times()*,
 26151 *ulimit()*, *umask()*

26152 XBD Chapter 8 (on page 173), **<unistd.h>**26153 **CHANGE HISTORY**

26154 First released in Issue 1. Derived from Issue 1 of the SVID.

26155 **Issue 5**26156 The DESCRIPTION is updated for alignment with the POSIX Realtime Extension and the POSIX
 26157 Threads Extension.

26158 Large File Summit extensions are added.

26159 **Issue 6**26160 The following new requirements on POSIX implementations derive from alignment with the
 26161 Single UNIX Specification:

- 26162 • In the DESCRIPTION, behavior is defined for when the process image file is not a valid
 26163 executable.
- 26164 • In this version, `_POSIX_SAVED_IDS` is mandated, thus the effective user ID and effective
 26165 group ID of the new process image shall be saved (as the saved set-user-ID and the saved
 26166 set-group-ID) for use by the *setuid()* function.
- 26167 • The [ELOOP] mandatory error condition is added.
- 26168 • A second [ENAMETOOLONG] is added as an optional error condition.
- 26169 • The [ETXTBSY] optional error condition is added.

26170 The following changes were made to align with the IEEE P1003.1a draft standard:

- 26171 • The [EINVAL] mandatory error condition is added.
- 26172 • The [ELOOP] optional error condition is added.

26173 The description of CPU-time clock semantics is added for alignment with IEEE Std 1003.1d-1999.

26174 The DESCRIPTION is updated for alignment with IEEE Std 1003.1j-2000 by adding semantics
 26175 for typed memory.

26176 The normative text is updated to avoid use of the term “must” for application requirements.

26177 The description of tracing semantics is added for alignment with IEEE Std 1003.1q-2000.

26178 IEEE PASC Interpretation 1003.1 #132 is applied.

26179 The DESCRIPTION is updated to make it explicit that the floating-point environment in the new
 26180 process image is set to the default.26181 The DESCRIPTION and RATIONALE are updated to include clarifications of how the contents
 26182 of a process image file affect the behavior of the *exec* functions.26183 IEEE Std 1003.1-2001/Cor 1-2002, item XSH/TC1/D6/15 is applied, adding a new paragraph to
 26184 the DESCRIPTION and text to the end of the APPLICATION USAGE section. This change

- 26185 addresses a security concern, where implementations may want to reopen file descriptors 0, 1,
26186 and 2 for programs with the set-user-id or set-group-id file mode bits calling the *exec* family of
26187 functions.
- 26188 IEEE Std 1003.1-2001/Cor 2-2004, item XSH/TC2/D6/24 is applied, applying changes to the
26189 DESCRIPTION, addressing which attributes are inherited by threads, and behavioral
26190 requirements for threads attributes.
- 26191 IEEE Std 1003.1-2001/Cor 2-2004, item XSH/TC2/D6/25 is applied, updating text in the
26192 RATIONALE from “the process signal mask be unchanged across an *exec*” to “the new process
26193 image inherits the signal mask of the thread that called *exec* in the old process image”.
- 26194 **Issue 7**
- 26195 Austin Group Interpretation 1003.1-2001 #047 is applied, adding the description of `_CS_V7_ENV`
26196 to the APPLICATION USAGE.
- 26197 Austin Group Interpretation 1003.1-2001 #143 is applied.
- 26198 The *fexecve()* function is added from The Open Group Technical Standard, 2006, Extended API
26199 Set Part 2.
- 26200 Functionality relating to the Asynchronous Input and Output, Memory Mapped Files, Threads,
26201 and Timers options is moved to the Base.
- 26202 Changes are made related to support for finegrained timestamps.