

68659 **NAME**68660 `wait, waitpid` — wait for a child process to stop or terminate68661 **SYNOPSIS**68662 `#include <sys/wait.h>`68663 `pid_t wait(int *stat_loc);`68664 `pid_t waitpid(pid_t pid, int *stat_loc, int options);`68665 **DESCRIPTION**

68666 The `wait()` and `waitpid()` functions shall obtain status information pertaining to one of the
 68667 caller's child processes. Various options permit status information to be obtained for child
 68668 processes that have terminated or stopped. If status information is available for two or more
 68669 child processes, the order in which their status is reported is unspecified.

68670 The `wait()` function shall suspend execution of the calling thread until status information for one
 68671 of the terminated child processes of the calling process is available, or until delivery of a signal
 68672 whose action is either to execute a signal-catching function or to terminate the process. If more
 68673 than one thread is suspended in `wait()` or `waitpid()` awaiting termination of the same process,
 68674 exactly one thread shall return the process status at the time of the target process termination. If
 68675 status information is available prior to the call to `wait()`, return shall be immediate.

68676 The `waitpid()` function shall be equivalent to `wait()` if the `pid` argument is `(pid_t)-1` and the
 68677 `options` argument is 0. Otherwise, its behavior shall be modified by the values of the `pid` and
 68678 `options` arguments.

68679 The `pid` argument specifies a set of child processes for which `status` is requested. The `waitpid()`
 68680 function shall only return the status of a child process from this set:

- 68681 • If `pid` is equal to `(pid_t)-1`, `status` is requested for any child process. In this respect,
 68682 `waitpid()` is then equivalent to `wait()`.
- 68683 • If `pid` is greater than 0, it specifies the process ID of a single child process for which `status` is
 68684 requested.
- 68685 • If `pid` is 0, `status` is requested for any child process whose process group ID is equal to that
 68686 of the calling process.
- 68687 • If `pid` is less than `(pid_t)-1`, `status` is requested for any child process whose process group
 68688 ID is equal to the absolute value of `pid`.

68689 The `options` argument is constructed from the bitwise-inclusive OR of zero or more of the
 68690 following flags, defined in the `<sys/wait.h>` header:

68691 XSI **WCONTINUED** The `waitpid()` function shall report the status of any continued child process
 68692 specified by `pid` whose status has not been reported since it continued from a
 68693 job control stop.

68694 **WNOHANG** The `waitpid()` function shall not suspend execution of the calling thread if
 68695 `status` is not immediately available for one of the child processes specified by
 68696 `pid`.

68697 **WUNTRACED** The status of any child processes specified by `pid` that are stopped, and whose
 68698 status has not yet been reported since they stopped, shall also be reported to
 68699 the requesting process.

68700 XSI If the calling process has `SA_NOCLDWAIT` set or has `SIGCHLD` set to `SIG_IGN`, and the process
 68701 has no unwaited-for children that were transformed into zombie processes, the calling thread
 68702 shall block until all of the children of the process containing the calling thread terminate, and
 68703 `wait()` and `waitpid()` shall fail and set `errno` to `[ECHILD]`.

If *wait()* or *waitpid()* return because the status of a child process is available, these functions shall return a value equal to the process ID of the child process. In this case, if the value of the argument *stat_loc* is not a null pointer, information shall be stored in the location pointed to by *stat_loc*. The value stored at the location pointed to by *stat_loc* shall be 0 if and only if the status returned is from a terminated child process that terminated by one of the following means:

1. The process returned 0 from *main()*.
2. The process called *_exit()* or *exit()* with a *status* argument of 0.
3. The process was terminated because the last thread in the process terminated.

Regardless of its value, this information may be interpreted using the following macros, which are defined in **<sys/wait.h>** and evaluate to integral expressions; the *stat_val* argument is the integer value pointed to by *stat_loc*.

WIFEXITED(*stat_val*)

Evaluates to a non-zero value if *status* was returned for a child process that terminated normally.

WEXITSTATUS(*stat_val*)

If the value of WIFEXITED(*stat_val*) is non-zero, this macro evaluates to the low-order 8 bits of the *status* argument that the child process passed to *_exit()* or *exit()*, or the value the child process returned from *main()*.

WIFSIGNALED(*stat_val*)

Evaluates to a non-zero value if *status* was returned for a child process that terminated due to the receipt of a signal that was not caught (see **<signal.h>**).

WTERMSIG(*stat_val*)

If the value of WIFSIGNALED(*stat_val*) is non-zero, this macro evaluates to the number of the signal that caused the termination of the child process.

WIFSTOPPED(*stat_val*)

Evaluates to a non-zero value if *status* was returned for a child process that is currently stopped.

WSTOPSIG(*stat_val*)

If the value of WIFSTOPPED(*stat_val*) is non-zero, this macro evaluates to the number of the signal that caused the child process to stop.

WIFCONTINUED(*stat_val*)

Evaluates to a non-zero value if *status* was returned for a child process that has continued from a job control stop.

It is unspecified whether the *status* value returned by calls to *wait()* or *waitpid()* for processes created by *posix_spawn()* or *posix_spawnnp()* can indicate a WIFSTOPPED(*stat_val*) before subsequent calls to *wait()* or *waitpid()* indicate WIFEXITED(*stat_val*) as the result of an error detected before the new process image starts executing.

It is unspecified whether the *status* value returned by calls to *wait()* or *waitpid()* for processes created by *posix_spawn()* or *posix_spawnnp()* can indicate a WIFSIGNALED(*stat_val*) if a signal is sent to the parent's process group after *posix_spawn()* or *posix_spawnnp()* is called.

If the information pointed to by *stat_loc* was stored by a call to *waitpid()* that specified the WUNTRACED flag and did not specify the WCONTINUED flag, exactly one of the macros WIFEXITED(**stat_loc*), WIFSIGNALED(**stat_loc*), and WIFSTOPPED(**stat_loc*) shall evaluate to a non-zero value.

68748 If the information pointed to by *stat_loc* was stored by a call to *waitpid()* that specified the
 68749 XSI WUNTRACED and WCONTINUED flags, exactly one of the macros WIFEXITED(**stat_loc*),
 68750 XSI WIFSIGNALED(**stat_loc*), WIFSTOPPED(**stat_loc*), and WIFCONTINUED(**stat_loc*) shall
 68751 evaluate to a non-zero value.

68752 If the information pointed to by *stat_loc* was stored by a call to *waitpid()* that did not specify the
 68753 XSI WUNTRACED or WCONTINUED flags, or by a call to the *wait()* function, exactly one of the
 68754 macros WIFEXITED(**stat_loc*) and WIFSIGNALED(**stat_loc*) shall evaluate to a non-zero value.

68755 If the information pointed to by *stat_loc* was stored by a call to *waitpid()* that did not specify the
 68756 XSI WUNTRACED flag and specified the WCONTINUED flag, or by a call to the *wait()* function,
 68757 XSI exactly one of the macros WIFEXITED(**stat_loc*), WIFSIGNALED(**stat_loc*), and
 68758 WIFCONTINUED(**stat_loc*) shall evaluate to a non-zero value.

68759 If _POSIX_REALTIME_SIGNALS is defined, and the implementation queues the SIGCHLD
 68760 signal, then if *wait()* or *waitpid()* returns because the status of a child process is available, any
 68761 pending SIGCHLD signal associated with the process ID of the child process shall be discarded.
 68762 Any other pending SIGCHLD signals shall remain pending.

68763 Otherwise, if SIGCHLD is blocked, if *wait()* or *waitpid()* return because the status of a child
 68764 process is available, any pending SIGCHLD signal shall be cleared unless the status of another
 68765 child process is available.

68766 For all other conditions, it is unspecified whether child *status* will be available when a SIGCHLD
 68767 signal is delivered.

68768 There may be additional implementation-defined circumstances under which *wait()* or *waitpid()*
 68769 report *status*. This shall not occur unless the calling process or one of its child processes
 68770 explicitly makes use of a non-standard extension. In these cases the interpretation of the
 68771 reported *status* is implementation-defined.

68772 If a parent process terminates without waiting for all of its child processes to terminate, the
 68773 remaining child processes shall be assigned a new parent process ID corresponding to an
 68774 implementation-defined system process.

68775 RETURN VALUE

68776 If *wait()* or *waitpid()* returns because the status of a child process is available, these functions
 68777 shall return a value equal to the process ID of the child process for which *status* is reported. If
 68778 *wait()* or *waitpid()* returns due to the delivery of a signal to the calling process, -1 shall be
 68779 returned and *errno* set to [EINTR]. If *waitpid()* was invoked with WNOHANG set in *options*, it
 68780 has at least one child process specified by *pid* for which *status* is not available, and *status* is not
 68781 available for any process specified by *pid*, 0 is returned. Otherwise, (*pid_t*)-1 shall be returned,
 68782 and *errno* set to indicate the error.

68783 ERRORS

68784 The *wait()* function shall fail if:

68785 [ECHILD] The calling process has no existing unwaited-for child processes.

68786 [EINTR] The function was interrupted by a signal. The value of the location pointed to
 68787 by *stat_loc* is undefined.

68788 The *waitpid()* function shall fail if:

68789 [ECHILD] The process specified by *pid* does not exist or is not a child of the calling
 68790 process, or the process group specified by *pid* does not exist or does not have
 68791 any member process that is a child of the calling process.

68792 [EINTR] The function was interrupted by a signal. The value of the location pointed to
 68793 by *stat_loc* is undefined.

68794 [EINVAL] The *options* argument is not valid.

68795 EXAMPLES

68796 Waiting for a Child Process and then Checking its Status

68797 The following example demonstrates the use of *waitpid()*, *fork()*, and the macros used to
 68798 interpret the status value returned by *waitpid()* (and *wait()*). The code segment creates a child
 68799 process which does some unspecified work. Meanwhile the parent loops performing calls to
 68800 *waitpid()* to monitor the status of the child. The loop terminates when child termination is
 68801 detected.

```

68802 #include <stdio.h>
68803 #include <stdlib.h>
68804 #include <unistd.h>
68805 #include <sys/wait.h>
68806 ...
68807 pid_t child_pid, wpid;
68808 int status;
68809
68810 child_pid = fork();
68811 if (child_pid == -1) {          /* fork() failed */
68812     perror("fork");
68813     exit(EXIT_FAILURE);
68814 }
68815 if (child_pid == 0) {          /* This is the child */
68816     /* Child does some work and then terminates */
68817     ...
68818 } else {                      /* This is the parent */
68819     do {
68820         wpid = waitpid(child_pid, &status, WUNTRACED
68821 #ifdef WCONTINUED          /* Not all implementations support this */
68822         | WCONTINUED
68823 #endif
68824         );
68825         if (wpid == -1) {
68826             perror("waitpid");
68827             exit(EXIT_FAILURE);
68828         }
68829         if (WIFEXITED(status)) {
68830             printf("child exited, status=%d\n", WEXITSTATUS(status));
68831         } else if (WIFSIGNALED(status)) {
68832             printf("child killed (signal %d)\n", WTERMSIG(status));
68833         } else if (WIFSTOPPED(status)) {
68834             printf("child stopped (signal %d)\n", WSTOPSIG(status));
68835 #ifdef WIFCONTINUED          /* Not all implementations support this */
68836         } else if (WIFCONTINUED(status)) {
68837             printf("child continued\n");

```

```

68837         #endif
68838             } else { /* Non-standard case -- may never happen */
68839                 printf("Unexpected status (0x%x)\n", status);
68840             }
68841         } while (!WIFEXITED(status) && !WIFSIGNALED(status));
68842     }

```

Waiting for a Child Process in a Signal Handler for SIGCHLD

The following example demonstrates how to use *waitpid()* in a signal handler for SIGCHLD without passing -1 as the *pid* argument. (See the APPLICATION USAGE section below for the reasons why passing a *pid* of -1 is not recommended.) The method used here relies on the standard behavior of *waitpid()* when SIGCHLD is blocked. On historical non-conforming systems, the status of some child processes might not be reported.

```

68849 #include <stdlib.h>
68850 #include <stdio.h>
68851 #include <signal.h>
68852 #include <sys/types.h>
68853 #include <sys/wait.h>
68854 #include <unistd.h>
68855 #define CHILDREN 10
68856 static void
68857 handle_sigchld(int signum, siginfo_t *sinfo, void *unused)
68858 {
68859     int status;
68860     /*
68861      * Obtain status information for the child which
68862      * caused the SIGCHLD signal and write its exit code
68863      * to stdout.
68864      */
68865     if (sinfo->si_code != CLD_EXITED)
68866     {
68867         static char msg[] = "wrong si_code\n";
68868         write(2, msg, sizeof msg - 1);
68869     }
68870     else if (waitpid(sinfo->si_pid, &status, 0) == -1)
68871     {
68872         static char msg[] = "waitpid() failed\n";
68873         write(2, msg, sizeof msg - 1);
68874     }
68875     else if (!WIFEXITED(status))
68876     {
68877         static char msg[] = "WIFEXITED was false\n";
68878         write(2, msg, sizeof msg - 1);
68879     }
68880     else
68881     {
68882         int code = WEXITSTATUS(status);
68883         char buf[2];
68884         buf[0] = '0' + code;

```

```

68885         buf[1] = '\n';
68886         write(1, buf, 2);
68887     }
68888 }
68889
68890 int
68891 main(void)
68892 {
68893     int i;
68894     pid_t pid;
68895     struct sigaction sa;
68896
68897     sa.sa_flags = SA_SIGINFO;
68898     sa.sa_sigaction = handle_sigchld;
68899     sigemptyset(&sa.sa_mask);
68900     if (sigaction(SIGCHLD, &sa, NULL) == -1)
68901     {
68902         perror("sigaction");
68903         exit(EXIT_FAILURE);
68904     }
68905
68906     for (i = 0; i < CHILDREN; i++)
68907     {
68908         switch (pid = fork())
68909         {
68910             case -1:
68911                 perror("fork");
68912                 exit(EXIT_FAILURE);
68913             case 0:
68914                 sleep(2);
68915                 _exit(i);
68916             }
68917         }
68918
68919         /* Wait for all the SIGCHLD signals, then terminate on SIGALRM */
68920         alarm(3);
68921         for (;;)
68922             pause();
68923     }
68924 }

```

APPLICATION USAGE

Calls to *wait()* will collect information about any child process. This may result in interactions with other interfaces that may be waiting for their own children (such as by use of *system()*). For this and other reasons it is recommended that portable applications not use *wait()*, but instead use *waitpid()*. For these same reasons, the use of *waitpid()* with a *pid* argument of *-1*, and the use of *waitid()* with the *idtype* argument set to *P_ALL*, are also not recommended for portable applications.

RATIONALE

A call to the *wait()* or *waitpid()* function only returns *status* on an immediate child process of the calling process; that is, a child that was produced by a single *fork()* call (perhaps followed by an *exec* or other function calls) from the parent. If a child produces grandchildren by further use of *fork()*, none of those grandchildren nor any of their descendants affect the behavior of a *wait()* from the original parent process. Nothing in this volume of POSIX.1-2008 prevents an implementation from providing extensions that permit a process to get *status* from a grandchild

or any other process, but a process that does not use such extensions must be guaranteed to see *status* from only its direct children.

The *waitpid()* function is provided for three reasons:

1. To support job control
2. To permit a non-blocking version of the *wait()* function
3. To permit a library routine, such as *system()* or *pclose()*, to wait for its children without interfering with other terminated children for which the process has not waited

The first two of these facilities are based on the *wait3()* function provided by 4.3 BSD. The function uses the *options* argument, which is equivalent to an argument to *wait3()*. The WUNTRACED flag is used only in conjunction with job control on systems supporting job control. Its name comes from 4.3 BSD and refers to the fact that there are two types of stopped processes in that implementation: processes being traced via the *ptrace()* debugging facility and (untraced) processes stopped by job control signals. Since *ptrace()* is not part of this volume of POSIX.1-2008, only the second type is relevant. The name WUNTRACED was retained because its usage is the same, even though the name is not intuitively meaningful in this context.

The third reason for the *waitpid()* function is to permit independent sections of a process to spawn and wait for children without interfering with each other. For example, the following problem occurs in developing a portable shell, or command interpreter:

```
stream = popen("/bin/true");
(void) system("sleep 100");
(void) pclose(stream);
```

On all historical implementations, the final *pclose()* fails to reap the *wait()* *status* of the *popen()*.

The status values are retrieved by macros, rather than given as specific bit encodings as they are in most historical implementations (and thus expected by existing programs). This was necessary to eliminate a limitation on the number of signals an implementation can support that was inherent in the traditional encodings. This volume of POSIX.1-2008 does require that a *status* value of zero corresponds to a process calling *_exit(0)*, as this is the most common encoding expected by existing programs. Some of the macro names were adopted from 4.3 BSD.

These macros syntactically operate on an arbitrary integer value. The behavior is undefined unless that value is one stored by a successful call to *wait()* or *waitpid()* in the location pointed to by the *stat_loc* argument. An early proposal attempted to make this clearer by specifying each argument as **stat_loc* rather than *stat_val*. However, that did not follow the conventions of other specifications in this volume of POSIX.1-2008 or traditional usage. It also could have implied that the argument to the macro must literally be **stat_loc*; in fact, that value can be stored or passed as an argument to other functions before being interpreted by these macros.

The extension that affects *wait()* and *waitpid()* and is common in historical implementations is the *ptrace()* function. It is called by a child process and causes that child to stop and return a *status* that appears identical to the *status* indicated by WIFSTOPPED. The *status* of *ptrace()* children is traditionally returned regardless of the WUNTRACED flag (or by the *wait()* function). Most applications do not need to concern themselves with such extensions because they have control over what extensions they or their children use. However, applications, such as command interpreters, that invoke arbitrary processes may see this behavior when those arbitrary processes misuse such extensions.

Implementations that support **core** file creation or other implementation-defined actions on termination of some processes traditionally provide a bit in the *status* returned by *wait()* to indicate that such actions have occurred.

Allowing the *wait()* family of functions to discard a pending SIGCHLD signal that is associated with a successfully waited-for child process puts them into the *sigwait()* and *sigwaitinfo()* category with respect to SIGCHLD.

This definition allows implementations to treat a pending SIGCHLD signal as accepted by the process in *wait()*, with the same meaning of “accepted” as when that word is applied to the *sigwait()* family of functions.

Allowing the *wait()* family of functions to behave this way permits an implementation to be able to deal precisely with SIGCHLD signals.

In particular, an implementation that does accept (discard) the SIGCHLD signal can make the following guarantees regardless of the queuing depth of signals in general (the list of waitable children can hold the SIGCHLD queue):

1. If a SIGCHLD signal handler is established via *sighandler()* without the SA_RESETHAND flag, SIGCHLD signals can be accurately counted; that is, exactly one SIGCHLD signal will be delivered to or accepted by the process for every child process that terminates.
2. A single *wait()* issued from a SIGCHLD signal handler can be guaranteed to return immediately with status information for a child process.
3. When SA_SIGINFO is requested, the SIGCHLD signal handler can be guaranteed to receive a non-null pointer to a **siginfo_t** structure that describes a child process for which a wait via *waitpid()* or *waitid()* will not block or fail.
4. The *system()* function will not cause the SIGCHLD handler of a process to be called as a result of the *fork()/exec* executed within *system()* because *system()* will accept the SIGCHLD signal when it performs a *waitpid()* for its child process. This is a desirable behavior of *system()* so that it can be used in a library without causing side-effects to the application linked with the library.

An implementation that does not permit the *wait()* family of functions to accept (discard) a pending SIGCHLD signal associated with a successfully waited-for child, cannot make the guarantees described above for the following reasons:

Guarantee #1

Although it might be assumed that reliable queuing of all SIGCHLD signals generated by the system can make this guarantee, the counter-example is the case of a process that blocks SIGCHLD and performs an indefinite loop of *fork()/wait()* operations. If the implementation supports queued signals, then eventually the system will run out of memory for the queue. The guarantee cannot be made because there must be some limit to the depth of queuing.

Guarantees #2 and #3

These cannot be guaranteed unless the *wait()* family of functions accepts the SIGCHLD signal. Otherwise, a *fork()/wait()* executed while SIGCHLD is blocked (as in the *system()* function) will result in an invocation of the handler when SIGCHLD is unblocked, after the process has disappeared.

Guarantee #4

Although possible to make this guarantee, *system()* would have to set the SIGCHLD handler to SIG_DFL so that the SIGCHLD signal generated by its *fork()* would be discarded (the SIGCHLD default action is to be ignored), then restore it to its previous setting. This would have the undesirable side-effect of discarding all SIGCHLD signals pending to the process.

69025 **FUTURE DIRECTIONS**

69026 None.

69027 **SEE ALSO**69028 *exec*, *exit()*, *fork()*, *system()*, *waitid()*69029 XBD Section 4.11 (on page 110), **<signal.h>**, **<sys/wait.h>**69030 **CHANGE HISTORY**

69031 First released in Issue 1. Derived from Issue 1 of the SVID.

69032 **Issue 5**

69033 The DESCRIPTION is updated for alignment with the POSIX Threads Extension.

69034 **Issue 6**69035 The following new requirements on POSIX implementations derive from alignment with the
69036 Single UNIX Specification:

- 69037
- The requirement to include **<sys/types.h>** has been removed. Although **<sys/types.h>** was
69038 required for conforming implementations of previous POSIX specifications, it was not
69039 required for UNIX applications.

69040 The following changes were made to align with the IEEE P1003.1a draft standard:

- 69041
- The processing of the SIGCHLD signal and the [ECHILD] error is clarified.

69042 The semantics of WIFSTOPPED(*stat_val*), WIFEXITED(*stat_val*), and WIFSIGNALED(*stat_val*)
69043 are defined with respect to *posix_spawn()* or *posix_spawnnp()* for alignment with IEEE Std
69044 1003.1d-1999.

69045 The DESCRIPTION is updated for alignment with the ISO/IEC 9899:1999 standard.

69046 IEEE Std 1003.1-2001/Cor 2-2004, item XSH/TC2/D6/145 is applied, adding the example to the
69047 EXAMPLES section.69048 **Issue 7**

69049 SD5-XSH-ERN-202 is applied.

69050 APPLICATION USAGE is added, recommending that the *wait()* function not be used.69051 An additional example for *waitpid()* is added.

NAME

waitid — wait for a child process to change state

SYNOPSIS

```
#include <sys/wait.h>
```

```
int waitid(idtype_t idtype, id_t id, siginfo_t *infop, int options);
```

DESCRIPTION

The *waitid()* function shall suspend the calling thread until one child of the process containing the calling thread changes state. It records the current state of a child in the structure pointed to by *infop*. The fields of the structure pointed to by *infop* are filled in as described for the SIGCHLD signal in <signal.h>. If a child process changed state prior to the call to *waitid()*, *waitid()* shall return immediately. If more than one thread is suspended in *wait()*, *waitid()*, or *waitpid()* waiting for termination of the same process, exactly one thread shall return the process status at the time of the target process termination.

The *idtype* and *id* arguments are used to specify which children *waitid()* waits for.

If *idtype* is P_PID, *waitid()* shall wait for the child with a process ID equal to (**pid_t**)*id*.

If *idtype* is P_PGID, *waitid()* shall wait for any child with a process group ID equal to (**pid_t**)*id*.

If *idtype* is P_ALL, *waitid()* shall wait for any children and *id* is ignored.

The *options* argument is used to specify which state changes *waitid()* shall wait for. It is formed by OR'ing together the following flags:

WCONTINUED Status shall be returned for any child that was stopped and has been continued.

WEXITED Wait for processes that have exited.

WNOHANG Do not hang if no status is available; return immediately.

WNOWAIT Keep the process whose status is returned in *infop* in a waitable state. This shall not affect the state of the process; the process may be waited for again after this call completes.

WSTOPPED Status shall be returned for any child that has stopped upon receipt of a signal.

Applications shall specify at least one of the flags WEXITED, WSTOPPED, or WCONTINUED to be OR'ed in with the *options* argument.

The application shall ensure that the *infop* argument points to a **siginfo_t** structure. If *waitid()* returns because a child process was found that satisfied the conditions indicated by the arguments *idtype* and *options*, then the structure pointed to by *infop* shall be filled in by the system with the status of the process. The *si_signo* member shall always be equal to SIGCHLD.

RETURN VALUE

If WNOHANG was specified and status is not available for any process specified by *idtype* and *id*, 0 shall be returned. If *waitid()* returns due to the change of state of one of its children, 0 shall be returned. Otherwise, -1 shall be returned and *errno* set to indicate the error.

ERRORS

The *waitid()* function shall fail if:

[ECHILD] The calling process has no existing unwaited-for child processes.

[EINTR] The *waitid()* function was interrupted by a signal.

69093 [EINVAL] An invalid value was specified for *options*, or *idtype* and *id* specify an invalid
 69094 set of processes.

69095 EXAMPLES

69096 None.

69097 APPLICATION USAGE

69098 Calls to *waitid()* with *idtype* equal to P_ALL will collect information about any child process.
 69099 This may result in interactions with other interfaces that may be waiting for their own children
 69100 (such as by use of *system()*). For this reason it is recommended that portable applications not
 69101 use *waitid()* with *idtype* of P_ALL. See also APPLICATION USAGE for *wait()*.

69102 RATIONALE

69103 None.

69104 FUTURE DIRECTIONS

69105 None.

69106 SEE ALSO

69107 *exec*, *exit()*, *wait()*

69108 XBD <signal.h>, <sys/wait.h>

69109 CHANGE HISTORY

69110 First released in Issue 4, Version 2.

69111 Issue 5

69112 Moved from X/OPEN UNIX extension to BASE.

69113 The DESCRIPTION is updated for alignment with the POSIX Threads Extension.

69114 Issue 6

69115 The normative text is updated to avoid use of the term “must” for application requirements.

69116 Issue 7

69117 Austin Group Interpretation 1003.1-2001 #060 is applied, updating the DESCRIPTION.

69118 The *waitid()* function is moved from the XSI option to the Base.

69119 APPLICATION USAGE is added, recommending that the *waitid()* function not be used with
 69120 *idtype* equal to P_ALL.

69121 The description of the WNOHANG flag is updated.

NAME

waitpid — wait for a child process to stop or terminate

SYNOPSIS

```
#include <sys/wait.h>
```

```
pid_t waitpid(pid_t pid, int *stat_loc, int options);
```

DESCRIPTION

Refer to *wait()*.